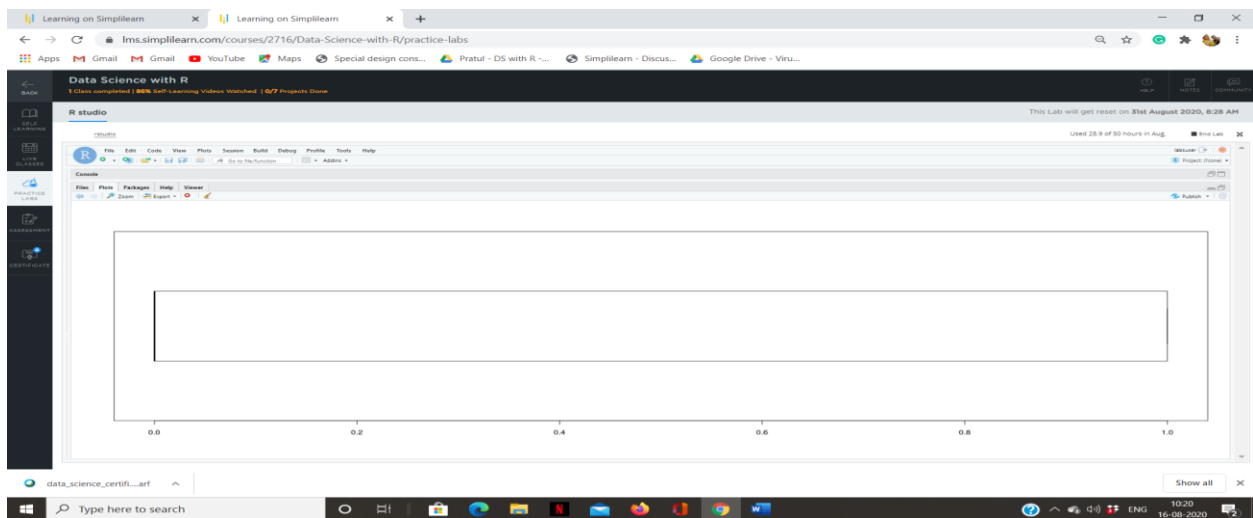


```

> #Analysis Tasks: Analyze the historical data and determine the key drivers for
admission.
> getwd()
[1] "/home/labsuser"
> dataset <- read.csv("College_admission.csv")
> View(dataset)
> #A->predictive tasks
> #task1 -to find missing value in data set.?if found perform missing value treat
ment.
> x=is.na(dataset)
> dataset[x]
numeric(0)
> #task2-Find outliers (if any, then perform outlier treatment)
> #although admit,gpa,ses,Gender_Male,Race,rank are cetagorical data still as per
said lets check it out for ouliers.
> boxplot(dataset$admit,horizontal = T)

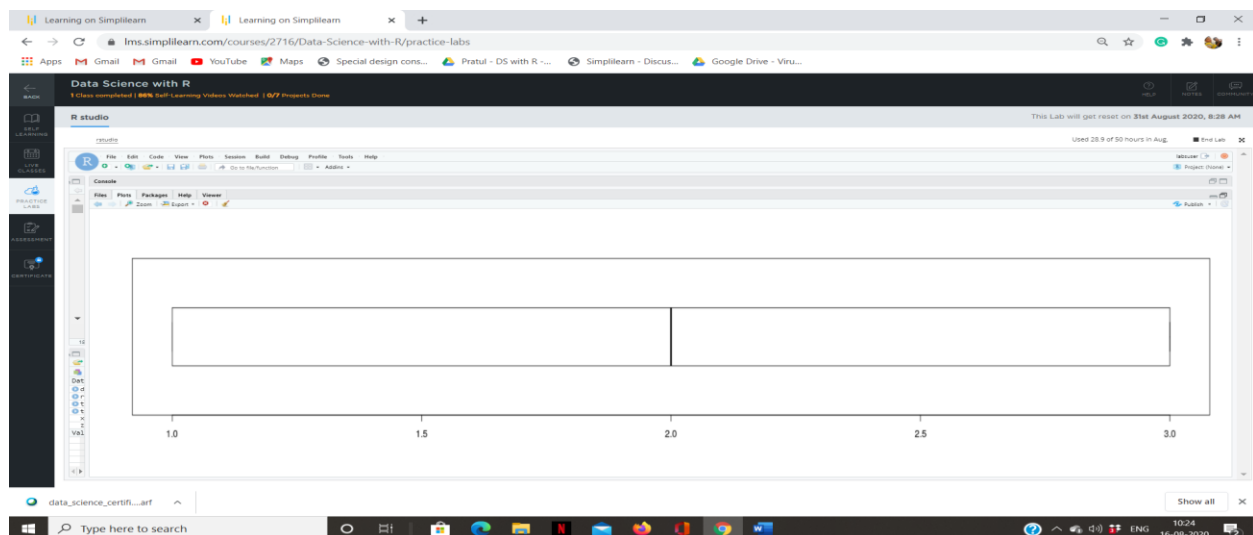
```



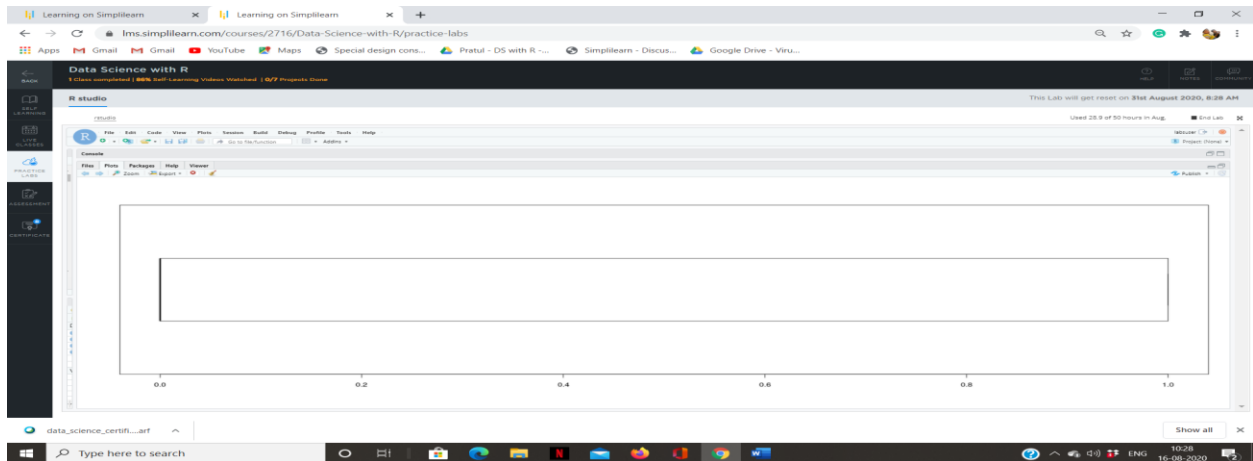
```

> boxplot(dataset$ses,horizontal = T)

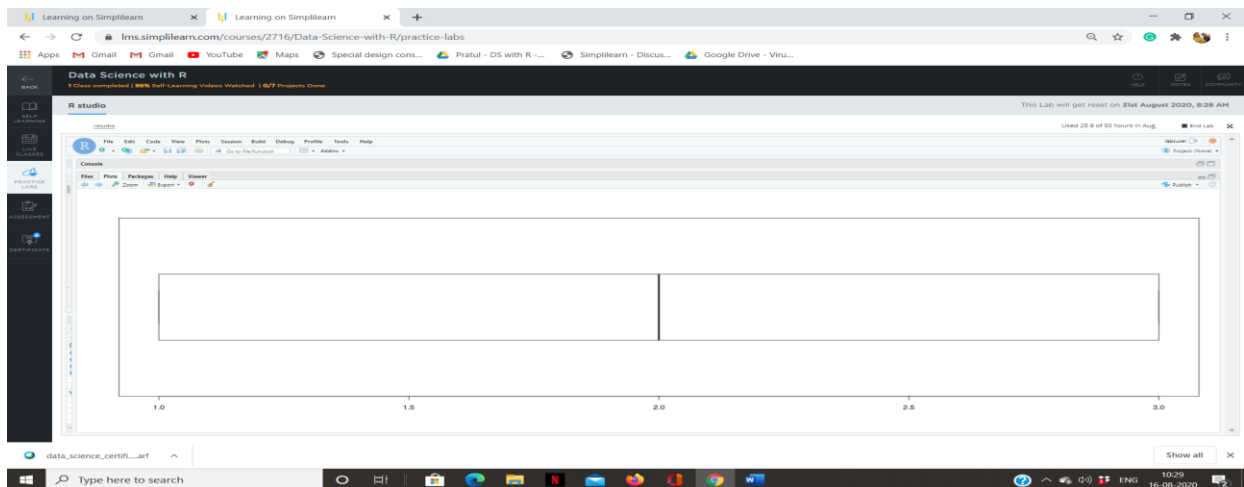
```



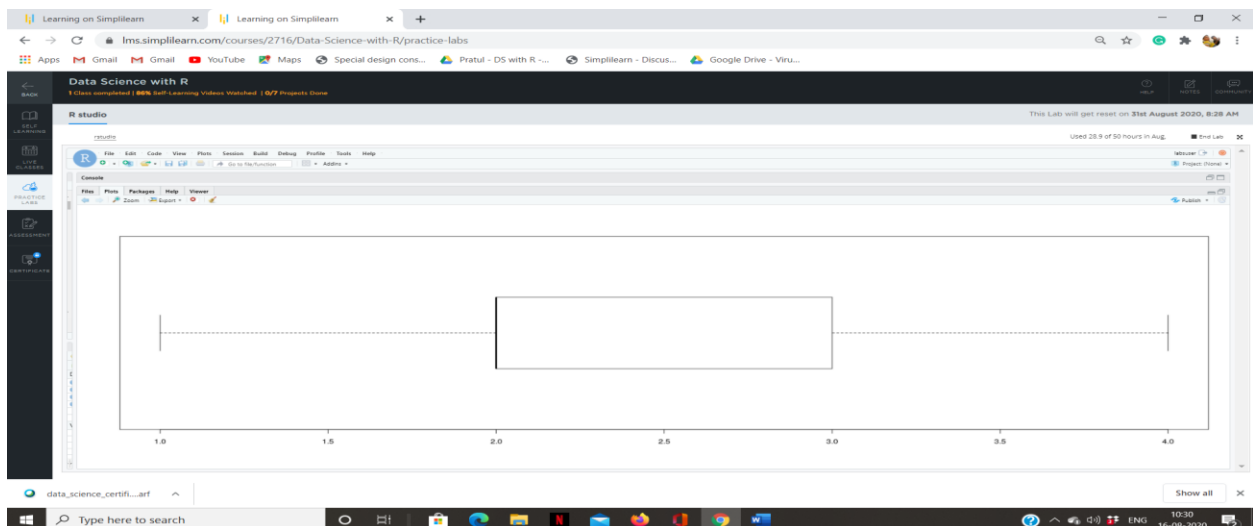
```
> boxplot(dataset$Gender_Male, horizontal = T)
```



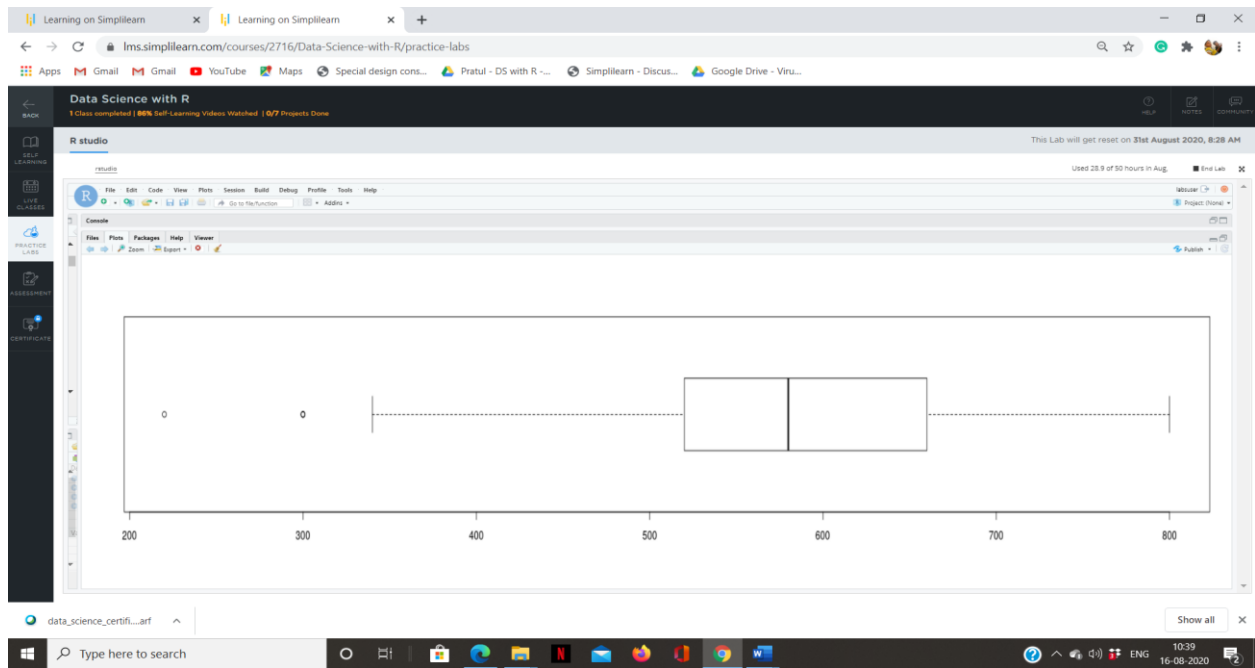
```
> boxplot(dataset$Race, horizontal = T)
```



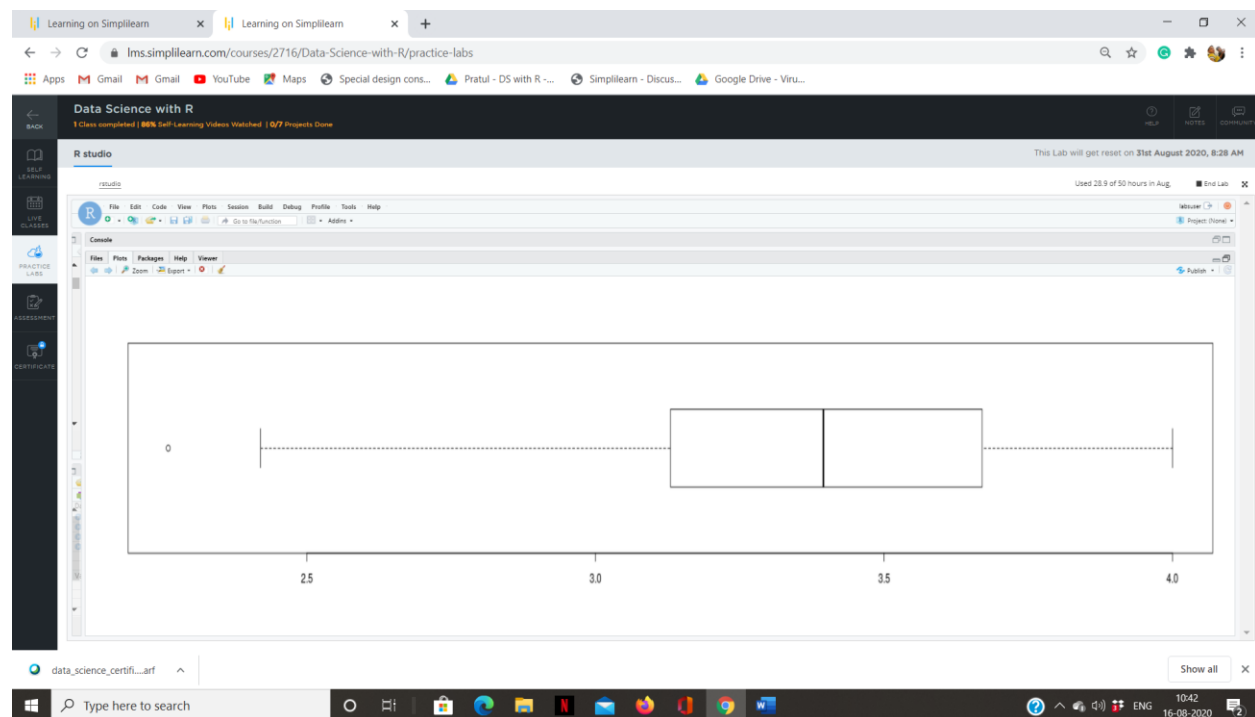
```
> boxplot(dataset$rank, horizontal = T)
```



```
OutVals = boxplot(dataset$gre, horizontal = T)$out
```



```
> OutVals
[1] 300 300 220 300
> #there are 4 outliers in dataset$gre or gre column
> outvals1=boxplot(dataset$gpa, horizontal = T)$out
```



```
> outvals1
[1] 2.26
```

```
> #task4-Find whether the data is normally distributed or not. Use the plot to determine the same.
```

```
> mean(dataset$gre)
```

```
[1] 590.777
```

```
> median(dataset$gre)
```

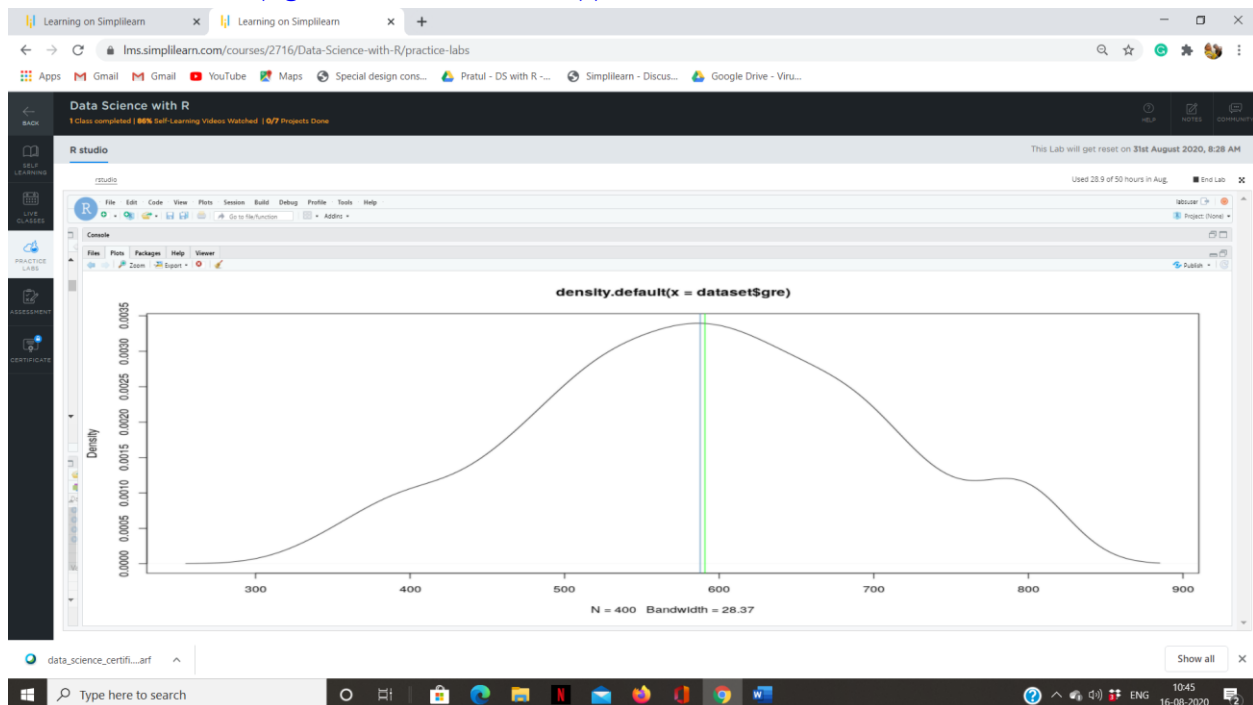
```
[1] 587.7
```

```
> #here mean is not equal to median in case of gre column of dataset hence data is not normally distributed
```

```
> #by visualization
```

```
> plot(density(dataset$gre))
```

```
> abline(v = c(mean(dataset$gre),  
+             median(dataset$gre)),  
+       col = c('green', 'steelblue'))
```



```
> #in visualization also mean is not equal to median in case of gre column of dataset hence data is not normally distributed
```

```
> mean(dataset$gpa)
```

```
[1] 3.392725
```

```
> median(dataset$gpa)
```

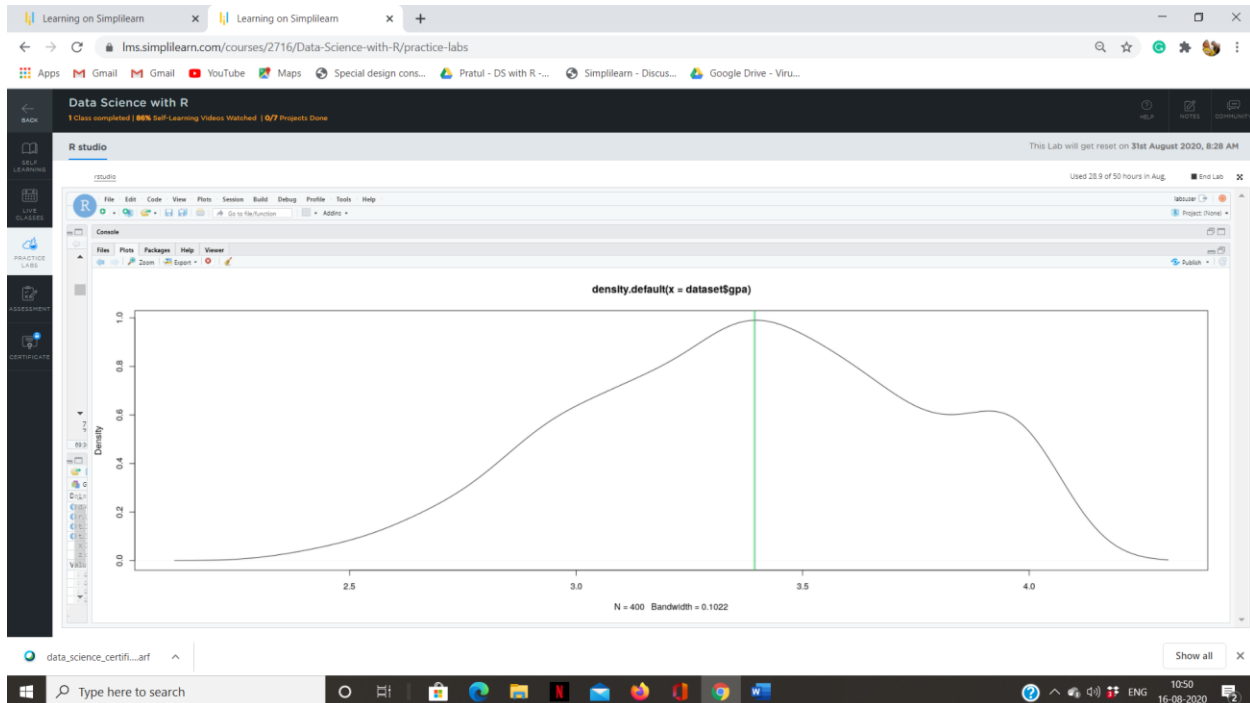
```
[1] 3.395
```

```
> #here mean is not equal to median in case of gpa column of dataset hence data is not normally distributed
```

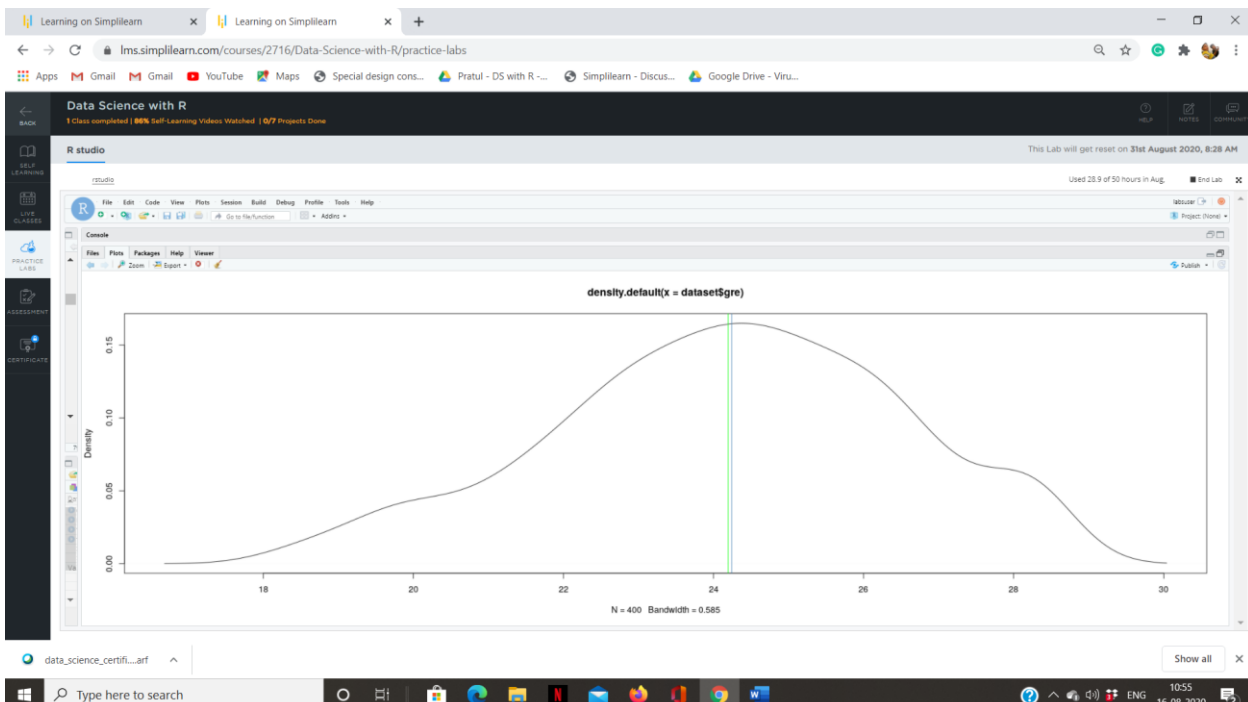
```
> #by visualization
```

```
> plot(density(dataset$gpa))
```

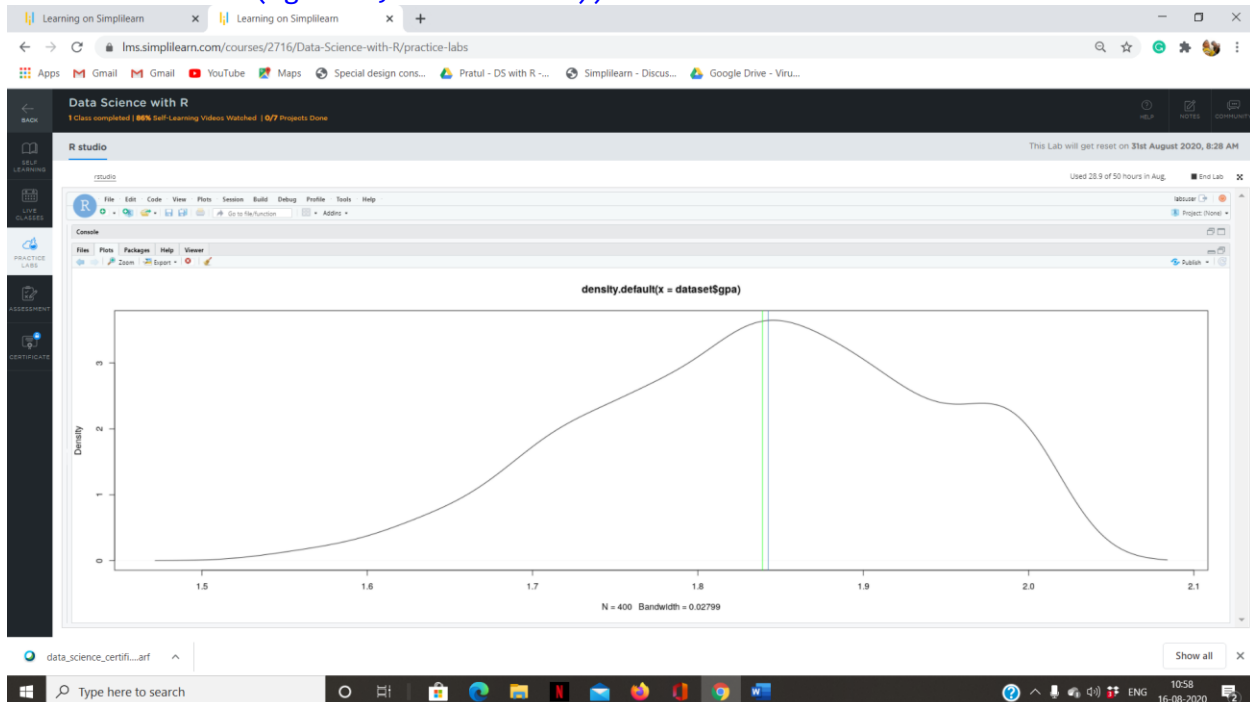
```
> abline(v = c(mean(dataset$gpa),  
+             median(dataset$gpa)),  
+       col = c('green', 'steelblue'))
```



```
> library(moments)
> skewness(dataset$gre, na.rm = TRUE)
[1] -0.02328672
> skewness(dataset$gpa, na.rm = TRUE)
[1] -0.1696461
> #both gre and gpa distribution have negative skewness.
> install.packages("ggpubr")
> library(magrittr)
> library(ggpubr)
> dataset$gre = sqrt(dataset$gre)
> dataset$gpa = sqrt(dataset$gpa)
> plot(density(dataset$gre))
> abline(v = c(mean(dataset$gre),
+             median(dataset$gre)),
+       col = c('green', 'steelblue'))
```



```
> plot(density(dataset$gpa))
> abline(v = c(mean(dataset$gpa),
+             median(dataset$gpa)),
+       col = c('green', 'steelblue'))
```



```
> mean(dataset$gre)
[1] 24.19508
> median(dataset$gre)
[1] 24.24252
> mean(dataset$gpa)
[1] 1.839056
> median(dataset$gpa)
[1] 1.842552
```

```
> #now the data is normally distributed
> #task6-Use variable reduction techniques to identify significant variables.
> # Splitting the dataset into the Training set and Test set
> install.packages('caTools')
> library(caTools)
> set.seed(125)
> split = sample.split(dataset$admit, SplitRatio = 0.8)
> training_set = subset(dataset, split == TRUE)
> test_set = subset(dataset, split == FALSE)
> ## Feature Scaling
> training_set[,2:3] = scale(training_set[,2:3])
> test_set[,2:3] = scale(test_set[,2:3])
> # Fitting Multiple Linear Regression to the Training set
> regressor = lm(formula = admit ~., data = training_set)
> summary(regressor)
```

```
Call:
lm(formula = admit ~ ., data = training_set)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.7908	-0.3375	-0.1726	0.4673	0.9247

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.619667	0.083188	7.449	9.47e-13 ***
gre	0.061700	0.026881	2.295	0.022385 *
gpa	0.048105	0.026803	1.795	0.073673 .
ses2	-0.009216	0.060221	-0.153	0.878463
ses3	-0.029756	0.061847	-0.481	0.630773
Gender_Male1	-0.059993	0.050035	-1.199	0.231439
Race2	-0.124931	0.060298	-2.072	0.039104 *
Race3	-0.065092	0.061066	-1.066	0.287286
rank2	-0.140065	0.075624	-1.852	0.064962 .
rank3	-0.299029	0.079183	-3.776	0.000191 ***
rank4	-0.345314	0.088168	-3.917	0.000111 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4422 on 309 degrees of freedom  
Multiple R-squared: 0.1304, Adjusted R-squared: 0.1023  
F-statistic: 4.634 on 10 and 309 DF, p-value: 3.678e-06

```
> #backward elimination for eliminating
> regressor = lm(formula = admit ~gre+gpa+rank+Race+ses+Gender_Male,data = training_set)
> summary(regressor)
```

Call:

```
lm(formula = admit ~ gre + gpa + rank + Race + ses + Gender_Male,
    data = training_set)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.7908	-0.3375	-0.1726	0.4673	0.9247

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.619667	0.083188	7.449	9.47e-13 ***
gre	0.061700	0.026881	2.295	0.022385 *
gpa	0.048105	0.026803	1.795	0.073673 .
rank2	-0.140065	0.075624	-1.852	0.064962 .
rank3	-0.299029	0.079183	-3.776	0.000191 ***
rank4	-0.345314	0.088168	-3.917	0.000111 ***
Race2	-0.124931	0.060298	-2.072	0.039104 *
Race3	-0.065092	0.061066	-1.066	0.287286
ses2	-0.009216	0.060221	-0.153	0.878463
ses3	-0.029756	0.061847	-0.481	0.630773

```
Gender_Male1 -0.059993 0.050035 -1.199 0.231439
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.4422 on 309 degrees of freedom
```

```
Multiple R-squared: 0.1304, Adjusted R-squared: 0.1023
```

```
F-statistic: 4.634 on 10 and 309 DF, p-value: 3.678e-06
```

```
> regressor = lm(formula = admit ~gre+gpa+rank+Race+ses,data = training_set)
> summary(regressor)
```

```
Call:
```

```
lm(formula = admit ~ gre + gpa + rank + Race + ses, data = training_set)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-0.7540	-0.3380	-0.1639	0.4571	0.9529

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	0.58534	0.07816	7.489	7.28e-13	***
gre	0.06102	0.02689	2.269	0.023950	*
gpa	0.04955	0.02680	1.849	0.065389	.
rank2	-0.13603	0.07560	-1.799	0.072950	.
rank3	-0.29130	0.07898	-3.688	0.000266	***
rank4	-0.34285	0.08821	-3.887	0.000124	***
Race2	-0.11880	0.06012	-1.976	0.049051	*
Race3	-0.06316	0.06109	-1.034	0.301985	
ses2	-0.01258	0.06020	-0.209	0.834602	
ses3	-0.02855	0.06188	-0.461	0.644872	

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.4425 on 310 degrees of freedom
```

```
Multiple R-squared: 0.1264, Adjusted R-squared: 0.101
```

```
F-statistic: 4.983 on 9 and 310 DF, p-value: 2.839e-06
```

```
> regressor = lm(formula = admit ~gre+gpa+rank+ses,data = training_set)
> summary(regressor)
```

```
Call:
```

```
lm(formula = admit ~ gre + gpa + rank + ses, data = training_set)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-0.6831	-0.3413	-0.1701	0.4849	0.9472

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	0.52598	0.07127	7.380	1.44e-12	***
gre	0.06191	0.02693	2.299	0.022155	*
gpa	0.04800	0.02679	1.792	0.074128	.



```
rank2      -0.12157      0.07522   -1.616  0.107052
rank3      -0.28682      0.07890   -3.635  0.000325 ***
rank4      -0.33988      0.08844   -3.843  0.000147 ***
ses2       -0.02341      0.05996   -0.390  0.696434
ses3       -0.03625      0.06168   -0.588  0.557135
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4439 on 312 degrees of freedom  
Multiple R-squared: 0.1153, Adjusted R-squared: 0.09545  
F-statistic: 5.809 on 7 and 312 DF, p-value: 2.408e-06

```
> regressor = lm(formula = admit ~gre+gpa+rank,data = training_set)
> summary(regressor)
```

Call:

```
lm(formula = admit ~ gre + gpa + rank, data = training_set)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-0.6880 -0.3348 -0.1719  0.4777  0.9666
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.50776    0.06356   7.988 2.62e-14 ***
gre          0.06195    0.02680   2.312 0.021440 *
gpa          0.04774    0.02670   1.788 0.074758 .
rank2       -0.12267    0.07493  -1.637 0.102624
rank3       -0.28872    0.07860  -3.673 0.000281 ***
rank4       -0.34111    0.08809  -3.872 0.000131 ***
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4427 on 314 degrees of freedom  
Multiple R-squared: 0.1143, Adjusted R-squared: 0.1002  
F-statistic: 8.103 on 5 and 314 DF, p-value: 3.306e-07

```
> regressor = lm(formula = admit ~gpa+rank,data = training_set)
> summary(regressor)
```

Call:

```
lm(formula = admit ~ gpa + rank, data = training_set)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-0.6282 -0.3312 -0.1832  0.5010  0.9303
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.51839    0.06383   8.121 1.06e-14 ***
gpa          0.06993    0.02509   2.787 0.005645 **
rank2       -0.12915    0.07539  -1.713 0.087701 .
```

```
rank3      -0.30661    0.07875  -3.893 0.000121 ***
rank4      -0.35841    0.08837  -4.056 6.31e-05 ***
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.4458 on 315 degrees of freedom
```

```
Multiple R-squared:  0.09921,    Adjusted R-squared:  0.08777
```

```
F-statistic: 8.673 on 4 and 315 DF,  p-value: 1.186e-06
```

```
> #task7-Run logistic model to determine the factors that influence the admission
  process of a student (Drop insignificant variables)
```

```
> #Encoding the target feature as factor
```

```
> dataset$admit = factor(dataset$admit,)
```

```
> class(dataset$admit)
```

```
[1] "factor"
```

```
> levels(dataset$admit)
```

```
[1] "0" "1"
```

```
> # Splitting the dataset into the Training set and Test set
```

```
> install.packages('caTools')
```

```
Error in install.packages : Updating loaded packages
```

```
> library(caTools)
```

```
> set.seed(169)
```

```
> split = sample.split(dataset$admit,SplitRatio = 0.75)
```

```
> training_set = subset(dataset,split == TRUE)
```

```
> test_set = subset(dataset,split == FALSE)
```

```
> ## Feature Scaling
```

```
> training_set[,2:3] =scale(training_set[,2:3])
```

```
> test_set[,2:3] =scale(test_set[,2:3])
```

```
> # Fitting Logistic Regression to the Training set
```

```
> classifier = glm(formula = admit ~.,
```

```
+               family = binomial,
```

```
+               data = training_set)
```

```
> summary(classifier)
```

```
Call:
```

```
glm(formula = admit ~ ., family = binomial, data = training_set)
```

```
Deviance Residuals:
```

```
      Min       1Q   Median       3Q      Max
-1.8154  -0.8470  -0.5981   1.0547   2.1786
```

```
Coefficients:
```

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    0.6269    0.4091   1.533 0.125381
gre             0.3978    0.1520   2.617 0.008874 **
gpa             0.2210    0.1481   1.492 0.135590
ses2            -0.2520    0.3240  -0.778 0.436850
ses3            -0.4869    0.3277  -1.486 0.137297
Gender_Male1    -0.2514    0.2713  -0.926 0.354210
Race2           -0.5256    0.3329  -1.579 0.114408
Race3           -0.3322    0.3215  -1.033 0.301466
rank2           -0.6246    0.3648  -1.712 0.086886 .
```

```
rank3      -1.2316      0.3943  -3.123 0.001787 **
rank4      -1.7735      0.4936  -3.593 0.000327 ***
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.6 on 299 degrees of freedom  
Residual deviance: 333.0 on 289 degrees of freedom  
AIC: 355

Number of Fisher Scoring iterations: 4

```
> # Predicting the Test set results
> prob_pred = predict(classifier, type = 'response', newdata = test_set[, -1])
> y_pred = ifelse(prob_pred > 0.5, 1, 0)
> #making confusion matrix for test set
> m = table(test_set[, 1], y_pred > 0.5)
> m
```

	FALSE	TRUE
0	64	4
1	27	5

```
> #checking accuracy of the model over test set (according to task 8)
> accuracy = (64+5)/(64+27+5+4)*100 #69
> classifier = glm(formula = admit ~ gre + gpa + rank + Race + ses + Gender_Male,
+                  family = binomial,
+                  data = training_set)
> summary(classifier)
```

Call:

```
glm(formula = admit ~ gre + gpa + rank + Race + ses + Gender_Male,
    family = binomial, data = training_set)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.8154	-0.8470	-0.5981	1.0547	2.1786

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.6269	0.4091	1.533	0.125381
gre	0.3978	0.1520	2.617	0.008874 **
gpa	0.2210	0.1481	1.492	0.135590
rank2	-0.6246	0.3648	-1.712	0.086886 .
rank3	-1.2316	0.3943	-3.123	0.001787 **
rank4	-1.7735	0.4936	-3.593	0.000327 ***
Race2	-0.5256	0.3329	-1.579	0.114408
Race3	-0.3322	0.3215	-1.033	0.301466
ses2	-0.2520	0.3240	-0.778	0.436850
ses3	-0.4869	0.3277	-1.486	0.137297
Gender_Male1	-0.2514	0.2713	-0.926	0.354210

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.6 on 299 degrees of freedom  
Residual deviance: 333.0 on 289 degrees of freedom  
AIC: 355

Number of Fisher Scoring iterations: 4

```
> classifier = glm(formula = admit ~gre+gpa+rank+Race+ses,  
+                  family = binomial,  
+                  data = training_set)  
> summary(classifier)
```

Call:

```
glm(formula = admit ~ gre + gpa + rank + Race + ses, family = binomial,  
    data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.8027	-0.8546	-0.5903	1.0647	2.2240

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.4760	0.3731	1.276	0.201998
gre	0.3940	0.1517	2.598	0.009379 **
gpa	0.2282	0.1478	1.544	0.122566
rank2	-0.6128	0.3632	-1.687	0.091550 .
rank3	-1.1959	0.3906	-3.062	0.002198 **
rank4	-1.7505	0.4919	-3.559	0.000372 ***
Race2	-0.4983	0.3306	-1.507	0.131781
Race3	-0.3182	0.3208	-0.992	0.321215
ses2	-0.2654	0.3227	-0.822	0.410839
ses3	-0.4740	0.3269	-1.450	0.147093

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.60 on 299 degrees of freedom  
Residual deviance: 333.86 on 290 degrees of freedom  
AIC: 353.86

Number of Fisher Scoring iterations: 4

```
> classifier = glm(formula = admit ~gre+gpa+rank,  
+                  family = binomial,  
+                  data = training_set)  
> classifier = glm(formula = admit ~gre+gpa+rank+Race,  
+                  family = binomial,  
+                  data = training_set)
```

```
> summary(classifier)
```

Call:

```
glm(formula = admit ~ gre + gpa + rank + Race, family = binomial,  
     data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.7666	-0.8674	-0.6088	1.1129	2.2986

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.2665	0.3405	0.783	0.433772
gre	0.3986	0.1503	2.651	0.008018 **
gpa	0.2194	0.1471	1.492	0.135796
rank2	-0.6333	0.3625	-1.747	0.080606 .
rank3	-1.2158	0.3892	-3.124	0.001785 **
rank4	-1.7357	0.4893	-3.547	0.000389 ***
Race2	-0.5377	0.3281	-1.639	0.101278
Race3	-0.3012	0.3195	-0.943	0.345697

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.60 on 299 degrees of freedom  
Residual deviance: 336.03 on 292 degrees of freedom  
AIC: 352.03

Number of Fisher Scoring iterations: 4

```
> summary(classifier)
```

Call:

```
glm(formula = admit ~ gre + gpa + rank + Race, family = binomial,  
     data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.7666	-0.8674	-0.6088	1.1129	2.2986

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.2665	0.3405	0.783	0.433772
gre	0.3986	0.1503	2.651	0.008018 **
gpa	0.2194	0.1471	1.492	0.135796
rank2	-0.6333	0.3625	-1.747	0.080606 .
rank3	-1.2158	0.3892	-3.124	0.001785 **
rank4	-1.7357	0.4893	-3.547	0.000389 ***
Race2	-0.5377	0.3281	-1.639	0.101278
Race3	-0.3012	0.3195	-0.943	0.345697

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.60 on 299 degrees of freedom  
Residual deviance: 336.03 on 292 degrees of freedom  
AIC: 352.03

Number of Fisher Scoring iterations: 4

```
> classifier = glm(formula = admit ~gpa+rank,  
+                  family = binomial,  
+                  data = training_set)  
> summary(classifier)
```

Call:

```
glm(formula = admit ~ gpa + rank, family = binomial, data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.4310	-0.8831	-0.6258	1.1694	2.1385

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.0295	0.2828	0.104	0.916898
gpa	0.3497	0.1356	2.578	0.009941 **
rank2	-0.5647	0.3487	-1.620	0.105317
rank3	-1.2696	0.3790	-3.350	0.000809 ***
rank4	-1.7255	0.4783	-3.608	0.000309 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.60 on 299 degrees of freedom  
Residual deviance: 346.15 on 295 degrees of freedom  
AIC: 356.15

Number of Fisher Scoring iterations: 4

```
> classifier1 = glm(formula = admit ~rank+gpa,  
+                  family = binomial,  
+                  data = training_set)  
> summary(classifier1)
```

Call:

```
glm(formula = admit ~ rank + gpa, family = binomial, data = training_set)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.4310	-0.8831	-0.6258	1.1694	2.1385

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.0295	0.2828	0.104	0.916898
rank2	-0.5647	0.3487	-1.620	0.105317
rank3	-1.2696	0.3790	-3.350	0.000809 ***
rank4	-1.7255	0.4783	-3.608	0.000309 ***
gpa	0.3497	0.1356	2.578	0.009941 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 374.60 on 299 degrees of freedom  
Residual deviance: 346.15 on 295 degrees of freedom  
AIC: 356.15

Number of Fisher Scoring iterations: 4

```
> #task 8-Calculate the accuracy of the model and run validation techniques
> # Predicting the Test set results
> prob_pred1 = predict(classifier1, type = 'response',newdata = test_set[,-1])
> y_pred1 = ifelse(prob_pred1>0.5 ,1 ,0)
> #making confusion matrix for test set
> m = table(test_set[,1],y_pred1>0.5)
> m
```

	FALSE	TRUE
0	66	2
1	26	6

```
> accuracy_of_logistic_regression = ((66+6)/((66+26+6+2))*100 #72
> #task 9-Try other modelling techniques like decision tree and SVM and select a
champion model
> #1.svm
> # Fitting SVM to the Training set, with kernel as linear
> install.packages(e1071)
> classifier = svm(formula = admit~.,
+                   data = training_set,
+                   type = 'C-classification',
+                   kernel = 'linear')
> library(e1071)
> # Predicting the Test set results
> y_pred = predict(classifier,newdata = test_set[,-1])
> # Making the Confusion Matrix
> cm = table(test_set[,1],y_pred)
> cm
```

	y_pred	
	0	1
0	68	0
1	32	0

```
> accuracy_of_svm= (68/100*100)
> classifier = svm(formula = admit~.,
+                   data = training_set,
```

```

+             type = 'C-classification',
+             kernel = 'linear')
> # Predicting the Test set results
> y_pred = predict(classifier,newdata = test_set[,-1])
> # Making the Confusion Matrix
> cm = table(test_set[,1],y_pred)
> cm
  y_pred
    0  1
0 68  0
1 32  0
> # Fitting SVM to the Training set, with kernel as radial
> classifier = svm(formula = admit~.,
+                 data = training_set,
+                 type = 'C-classification',
+                 kernel = 'radial')
> # Predicting the Test set results
> y_pred = predict(classifier,newdata = test_set[,-1])
> # Making the Confusion Matrix
> cm = table(test_set[,1],y_pred)
> cm
  y_pred
    0  1
0 68  0
1 32  0
> accuracy =68/100*100 #68
> #2.knn
> # Fitting KNN to the Training set
> install.packages(class)
> y_pred = knn(train = training_set[,-1],
+             test = test_set[,-1],
+             cl = training_set[,1],
+             k = 25,
+             prob = T) #3,5,7,9
> library(class)
> # Making the Confusion Matrix
> cm = table(test_set[,1],y_pred)
> cm
  y_pred
    0  1
0 62  6
1 26  6
> accuracy_of_knn = (62+6)/(6+62+6+26)*100 #68
> #3-decision tree
> # Fitting Decision Tree Classification to the Training set
> # Fitting Decision Tree Classification to the Training set
> install.packages('rpart')
> library(rpart)
> classifier = rpart(formula = admit~ .,
+                 data = training_set)
> # Predicting the Test set results
> y_pred = predict(classifier, newdata = test_set[-1],type = 'class')

```



```

> # Making the Confusion Matrix
> cm = table(test_set[,1], y_pred)
> cm
  y_pred
    0    1
0 58 10
1 22 10
> accuracy_of_decision_tree = (58+10)/(58+10+22+14)*100 #68

> #task10-Determine the accuracy rates for each kind of model.
> accuracy_of_logistic_regression = ((66+6)/(66+26+6+2))*100 #72
> accuracy_of_svm= (68+0)/(68+32+0+0)*100 #68%
> accuracy_of_knn = (62+6)/(6+62+6+26)*100 #68
> accuracy_of_decision_tree = (54+10)/(54+10+22+14)*100 #64
> #hance we chose logistic regression model and hance it is the chempion model.
> #task10-Determine the accuracy rates for each kind of model.
> accuracy_of_logistic_regression = ((66+6)/(66+26+6+2))*100 #72
> accuracy_of_svm= (68+0)/(68+32+0+0)*100 #68%
> accuracy_of_knn = (62+6)/(6+62+6+26)*100 #68
> accuracy_of_decision_tree = (54+10)/(54+10+22+14)*100 #64
> #hance we chose logistic regression model and hance it is the chempion model.
> # Importing the dataset
> # Fitting Random Forest Classification to the Training set
> install.packages('randomForest')
> library(randomForest)
randomForest 4.6-14
Type rfNews() to see new features/changes/bug fixes.

```

Attaching package: 'randomForest'

The following object is masked from 'package:ggplot2':

margin

```

> set.seed(321)
> classifier = randomForest(x = training_set[-1],
+                           y = training_set$admit,
+                           ntree = 180)
> # Predicting the Test set results
> y_pred = predict(classifier, newdata = test_set[-1])
> # Making the Confusion Matrix
> cm = table(test_set[, 1], y_pred)
> cm
  y_pred
    0    1
0 55 13
1 24   8

```

#B->Descriptive:

```

> #Categorize the average of grade point into High, Medium, and Low (with admission probability)
  a point chart.Cross grid for admission variables with GRE Categorization is shown below:

```

```

> #GRE Categorized
> #0-440 Low
> #440-580Medium
> #580+ High
> dataset <- read.csv("College_admission.csv")
> gre_grade1 = dataset$gre <= 440
> gre_grade2 = dataset$gre > 440 & dataset$gre <=580
> gre_grade3 = dataset$gre > 580
> dataset$gre[gre_grade1] = "low"
> dataset$gre[gre_grade2] = "medium"
> dataset$gre[gre_grade3] = "high"
> #Encoding the target feature as factor
> # Splitting the dataset into the Training set and Test set
> install.packages('caTools')
> library(caTools)
> set.seed(169)
> split = sample.split(dataset$admit,SplitRatio = 0.75)
> training_set = subset(dataset,split == TRUE)
> test_set = subset(dataset,split == FALSE)
> #now there can be two possible outcome of the task.
> #1>to compaire admission withe the gre column and making the visualization
> # Fitting Regression to the dataset
> regressor = lm(formula = admit ~ gre,data =dataset)
> summary(regressor)

```

Call:

```
lm(formula = admit ~ gre, data = dataset)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.3959	-0.3959	-0.2774	0.6041	0.8750

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.39594	0.03266	12.122	< 2e-16 ***
grelow	-0.27094	0.07379	-3.672	0.000274 ***
gremedium	-0.11852	0.04922	-2.408	0.016501 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4584 on 397 degrees of freedom

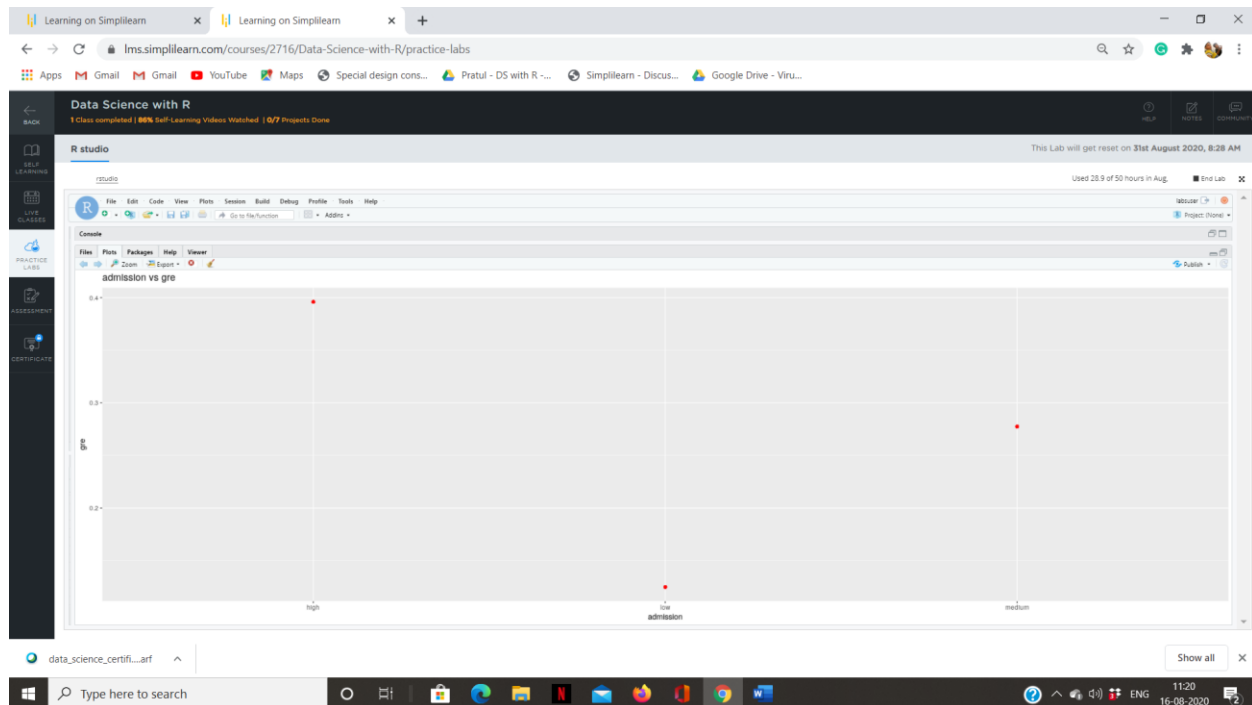
Multiple R-squared: 0.03738, Adjusted R-squared: 0.03253

F-statistic: 7.708 on 2 and 397 DF, p-value: 0.00052

```

> y =predict(regressor, type = 'response')
> library(ggplot2)
> ggplot() +
+   geom_point(aes(x=dataset$gre, y =predict(regressor, type = 'response')),
+             colour = 'red') +
+   ggtitle('admission vs gre ') +
+   xlab('admission')+
+   ylab('gre')

```



```
> #2>to compare the admission with all the other column and than making the visualization.
> regressor = lm(formula = admit ~ .,data =dataset)
> y =predict(regressor, type = 'response')
> summary(regressor)
```

Call:

```
lm(formula = admit ~ ., data = dataset)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.6606	-0.3434	-0.1892	0.4977	0.9811

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.24330	0.24206	1.005	0.3155
grelow	-0.16446	0.07675	-2.143	0.0327 *
gremedium	-0.07140	0.04977	-1.435	0.1522
gpa	0.15397	0.06324	2.435	0.0153 *
ses	-0.02667	0.02765	-0.965	0.3353
Gender_Male	-0.03037	0.04465	-0.680	0.4968
Race	-0.03175	0.02740	-1.159	0.2473
rank	-0.10884	0.02376	-4.581	6.22e-06 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4449 on 392 degrees of freedom

Multiple R-squared: 0.1049, Adjusted R-squared: 0.08892

F-statistic: 6.563 on 7 and 392 DF, p-value: 2.423e-07

```
> library(ggplot2)
> ggplot() +
+   geom_point(aes(x=dataset$gre, y =predict(regressor, type = 'response')),
+             colour = 'red') +
+   ggtitle('admission vs gre ') +
+   xlab('admission')+
+   ylab('gre')
```

